



FIELD GUIDE · FROM THE DAILY BRIEF

The AI UGC Playbook

Build an AI creator that posts for you — every day, fully automated, for under \$50 a month.

Compiled by Flowi Editorial · May 2026



ESSAY 01 · AI BUILDER

AI UGC: the unfair advantage every brand has in 2026

For two years UGC meant hiring creators and waiting weeks. One person with Claude and Higgsfield now out-publishes a five-creator roster for under \$50 a month. Here is the system, step by step.

Flowi Editorial

For two years, "UGC" — user-generated content — meant the same expensive loop. Find a creator. Negotiate a rate. Ship them product. Wait two weeks. Get three videos back. Hope one performs. Repeat.

That loop is over. Not because creators stopped working — because one person with Claude and Higgsfield can now produce more on-brand content in an afternoon than a five-creator roster produces in a month, for under \$50. This playbook is the exact system: no theory, no hype, the build step by step.

What you are actually building

You are not "making AI videos." You are building a content engine with four parts that run on a loop:

1. **An avatar** — a consistent face that fronts the content. A new AI influencer, or a clone of you.
2. **A persona** — who that avatar is: niche, voice, worldview, the things she posts about.
3. **A pipeline** — Claude writes the scripts and the prompts; Higgsfield generates the images and video.
4. **A cadence** — the loop runs daily; you review and post.

Once those four are set up, the marginal cost of the next piece of content is near zero. That is the whole unlock. Brands that understand it will out-publish brands that don't — not by a little, by an order of magnitude.

The economics, plainly

A single UGC creator runs \$150–500 per video. A modest roster is a few thousand dollars a month before you have posted anything.

The AI version is a Higgsfield subscription and a Claude subscription — call it \$40–50 a month, all in. The output is not "cheaper, lower-quality content." It is *more* content, *more* consistent, available the morning you need it instead of two weeks later. That cost gap is the unfair advantage — and it is temporary. The brands that move in 2026 get it cheap. By 2028 it is table stakes.

What this is not

Three honest caveats, because a playbook that oversells is worthless.

It is not passive. The system removes filming, scheduling, and the production grind. It does not remove judgment. Your job shifts from *making* content to *directing* it — choosing what is good, killing what is flat, steering the persona. Smaller job. Not zero.

It is not set-and-forget. Models update. A consistency trick that works this month may need adjusting next month. Treat this as a system you maintain, not a machine you walk away from.

It is not exempt from the rules. Every major platform now has disclosure expectations for AI-generated content. Label it where the platform asks. Audiences are fine with AI content that is honest about being AI content — they punish content that pretends. Done well, "this is an AI creator" is itself the hook.

How to read the rest of this playbook

The next four chapters are the build, in order:

- **Chapter 2 — the avatar.** New influencer or clone, and the one technical thing that makes or breaks everything: a face that never drifts.
- **Chapter 3 — the persona.** Turning a face into a creator with a niche, a voice and a reason to follow — using Claude as the strategist.
- **Chapter 4 — the content loop.** The daily pipeline: Claude writes, Higgsfield generates, with the exact prompts.
- **Chapter 5 — ship, scale, monetize.** Cadence, batching, running more than one avatar, and the honest paths to making money from this.

Work through them in order the first time. After that, this is a reference — come back to the chapter you need.

One last thing before you build. The creators who win with this system are not the ones with the best prompts. They are the ones with the best taste — who know what is worth posting. The tools are commodity now. Judgment is the moat. Keep that in mind through every step that follows.



ESSAY 02 · AI BUILDER

Building your AI avatar: a new influencer, or a clone of you

Every AI content system starts with one decision and one technical problem. Get the face right and everything downstream works. Here is how to choose your avatar and lock its identity so it never drifts.

Flowi Editorial

Every AI content system starts with one decision and one technical problem.

The decision: who is on camera? The problem: how do you keep that person looking like the same person across hundreds of posts? Get both right and everything downstream works. Get the second one wrong and the whole thing collapses — an audience will not follow a "person" whose face changes every video.

The decision: new influencer, or a clone of you

A brand-new AI influencer. You invent the person — age, look, name, whole identity. Full creative control, no personal exposure; your own face is never involved. Best when the brand is not *you*, when you want to run several avatars, or when you simply don't want to be on camera. The cost: you are building trust from zero. A new face has no history.

A clone of yourself. You train the avatar on your own face. The advantage is large and underrated — instant trust. An audience that already knows you does not care that a clip was AI-generated; it is *you*. Best when you have any existing audience, or when the brand is personal. The cost: it is your face. You inherit its upside and its risk.

There is no universally right answer. A practical default: if you have an audience, clone yourself; if you are starting clean or want to scale past one identity, build a new influencer. You can also do both — clone yourself for the main account, spin up new avatars for niche spin-offs.

The problem that breaks most attempts: drift

Here is why most people try AI UGC, post four times, and quit.

They generate each image fresh from a text prompt. Image one: a woman, late twenties, brown hair. Image two: also a woman, late twenties, brown hair — but not the *same* woman. The jaw is different, the eyes are different. Across a feed it reads as uncanny, and an audience will not attach to it.

This is *drift*. It is not solved with a better text prompt. It is solved with a trained character.

The fix: a Soul Character

Higgsfield's Soul Character exists for exactly this. You train it once, from reference images, and it produces a locked identity. Every generation after that is the *same* person — same face, same bone structure — placed in whatever scene you describe.

This is the single most important technical step in the playbook. Do not skip it and do not cut corners on it.

The reference image. A Soul Character is only as good as what you train it on. You need at least one clean, high-quality reference portrait:

- Sharp focus, even lighting, no heavy shadows across the face
- A neutral, natural expression — not an extreme one
- The face clearly visible, front-on or a gentle three-quarter angle
- Resolution high enough that skin texture is visible, not smoothed to plastic

Where it comes from depends on your path. Cloning yourself: take the photo properly — good light, plain background. Building a new influencer: generate several candidate portraits first, pick the strongest, and treat *that* as your reference. If you can supply several angles of the same face, the lock is tighter still.

The quality bar

When you generate your first batch with the trained character, judge it on four things, in this order:

1. **Identity.** Is it unmistakably the same person every time? Non-negotiable. Everything else is cosmetic.
2. **Realism.** Skin has texture and pores; AI tends to over-smooth. If it looks like a wax figure, it reads as fake — push for natural skin.
3. **Lighting consistency.** Light should behave the same way scene to scene. Wild lighting swings break the illusion as badly as a changing face.
4. **Hands and edges.** Still the classic tell. Check hands, jewelry, and where hair meets the background. Reject the failures — you have unlimited retries.

You are allowed to be picky. Generating another option costs cents; a weak avatar costs you the whole channel.

Once you have a locked, realistic, consistent face you are happy with, you have an avatar. In the next chapter you give her a brain — a persona that decides what she actually posts.



ESSAY 03 · AI BUILDER

The persona engine: making Claude define your AI creator

A face is not a creator. A creator is a face plus a persona — a niche, a voice, a worldview. Here is how to use Claude as your brand strategist to build one people actually follow.

Flowi Editorial

You now have a face that holds. A face is not a creator.

A creator is a face plus a *persona* — a niche, an audience, a voice, a worldview, a reason for anyone to follow. Skip this and you get what most AI accounts get: technically impressive images nobody follows, because there is no person there. The persona is what turns generated content into a channel. This is the step where Claude does the heavy lifting.

What a persona actually contains

A complete persona has five parts. Be vague on any of them and the content downstream will be vague too.

- **The niche.** Narrow. Not "fitness" — "strength training for women over 40." Not "AI" — "AI tools for solo real-estate agents." A narrow niche is easier to rank in, easier to monetize, and easier for Claude to write sharply for.
- **The audience.** One specific person. Their age, their job, what they want, what they are afraid of. You write for that one person, every post.
- **The voice.** Five words, no more. "Calm, precise, a little contrarian." The voice is the constraint that keeps a hundred posts feeling like one person.
- **The content pillars.** Three to five recurring themes. Every post belongs to one. Pillars are what stop you staring at a blank screen — you are never "thinking of content," you are filling a pillar.
- **The worldview.** The one belief the creator keeps returning to — the thing she is a little evangelical about. This is what makes a persona feel like a person and not a content dispenser.

The persona-brief prompt

Open Claude and give it this. Fill the bracket, then let it work:

"Act as a brand strategist. I am building an AI content creator in the [your niche] space.

Define a complete persona:

1. The exact target audience — one specific person, their situation, their goal, their fear.
2. A personality and tone of voice in exactly five words.
3. Five content pillars she posts on, each with a one-line description.

4. Her core worldview — the one belief she returns to.
5. Three things she would never say or do, to keep her consistent.
Be specific and slightly contrarian. Avoid anything generic enough to fit ten other creators."

That last instruction matters. Claude, unsteered, drifts toward the safe average. "Specific and slightly contrarian" pushes it off the average into something that can actually stand out.

Pressure-test before you build on it

Do not accept the first persona just because it reads well. Test it against three questions.

Is the niche monetizable? Does this audience buy things — courses, tools, products, services? "Aesthetic morning routines" gets views and sells nothing. "Tax strategy for freelancers" gets fewer views and sells everything. Want the second kind.

Is the voice distinct? Read the five voice words. Could they describe a hundred other accounts? If so, push back — ask Claude to "make the voice sharper, give it an edge most creators are too cautious to have."

Would you follow her? The honest gut check. If the persona does not interest *you*, it will not interest anyone. Iterate with Claude until it does — this is a conversation, not a one-shot prompt.

The persona document is your system brief

When the persona is right, save it into a single document — the persona doc.

That document is now the brief for everything. Every prompt in the next chapter begins with "using the persona below." The scripts reference it, the captions reference it, the visual direction references it. It is the single source of truth that keeps a hundred pieces of content — generated across a hundred different days — feeling like one coherent creator.

A face gives the channel a body. The persona gives it a mind. With both locked, you are ready to build the engine that produces the content — the next chapter.



ESSAY 04 · AI BUILDER

The content loop: Claude writes, Higgsfield generates

The engine that turns an avatar and a persona into daily content without filming anything. The exact pipeline, the content-plan prompt, and how to write prompts that render clean.

Flowi Editorial

You have a face that holds and a persona that gives it a mind. Now the engine — the loop that turns those two assets into content, every day, without you filming anything.

The loop has two halves. Claude is the writers' room. Higgsfield is the studio. Your job is the desk in between: you brief, you review, you post.

The pipeline, end to end

One full cycle:

1. **Claude plans and writes.** Working from the persona doc, it produces the content plan — hooks, scripts, captions — and, critically, the *image and video prompts* for each piece.
2. **Higgsfield generates.** You take Claude's visual prompts into Higgsfield, paired with your Soul Character, and it renders the actual images and clips — the same face, in the scene Claude described.
3. **You review.** Approve, tweak, or reject. Rejections cost cents; bad posts cost reach.
4. **You post.** Or schedule. The cycle resets.

The leverage is that steps 1 and 2 scale almost for free. Briefing Claude for seven days of content takes the same effort as briefing it for one.

The content-plan prompt

This is the workhorse. Give Claude the persona doc, then:

"Using the persona above, plan 7 days of short-form content. For each day give me:

- the content pillar it belongs to
- a scroll-stopping hook (the first line / first two seconds)
- a 20–30 second script in her voice
- the caption, written to her tone
- a detailed image-generation prompt for the visual — describe the scene, framing, lighting and mood; do NOT describe her face or identity

Optimise for watch time, shares and saves. Vary the hooks — never reuse a hook pattern twice."

Note the instruction on the image-prompt line: *do not describe her face or identity*. That is deliberate — and skipping it is the most common mistake.

How to write prompts Higgsfield renders well

The image prompt describes the **scene**. The Soul Character supplies the **person**. Keep those jobs separate.

A weak prompt: "a beautiful young woman with brown hair in a kitchen." You are now fighting your own Soul Character — re-describing a face the trained character already defines, which invites drift.

A strong prompt: "a bright modern kitchen, soft morning light through a window, she is mid-sentence holding a coffee mug, warm and candid, shot on 35mm, shallow depth of field." Scene, light, action, mood, camera. The character handles who she is; the prompt handles where she is and what is happening.

Three rules that hold up:

- **Always generate with the Soul Character attached.** Every single time. The consistency guarantee from Chapter 2 only works if you use it on every render.
- **Describe light and camera, not just objects.** "Soft window light, 35mm, shallow depth of field" is the difference between a render that looks like content and one that looks like a stock-photo generator.
- **One idea per image.** Cramming five elements into a prompt produces mush. Simple scene, executed cleanly.

The review gate is not optional

Once the loop runs, it is tempting to auto-post. Do not.

AI generation has a failure rate. A hand reads wrong. The lighting jumps. An expression lands as uncanny. One bad image in a feed of good ones does disproportionate damage — it is the post that makes someone think "this is fake" about *everything* you have made.

So there is one human gate: you, looking at every render before it goes out. It takes minutes a day. It is the cheapest insurance in the system — and it is the part "fully automated" marketing quietly skips. Run the loop automated; keep the gate human.

That is the engine. Claude briefs, Higgsfield renders, you approve. In the final chapter: how to run this at volume, across more than one avatar, and how it actually makes money.



ESSAY 05 · AI BUILDER

Ship, scale and monetize your AI UGC system

Turning a working content loop into a posting machine and then into money — cadence, batching, running multiple avatars, and the four honest paths to revenue.

Flowi Editorial

You have the engine. This chapter is about running it like an operator — turning a working loop into a posting machine, then into money.

Cadence: consistency beats volume

Pick a cadence you can hold, and hold it. One post a day, every day, beats seven in a burst and then silence. Algorithms reward reliability; so do audiences. A realistic start: one piece a day per account. That is modest, completely sustainable with this system — and already more than most human creators manage.

Batch — do not drip

The mistake is running the loop daily, in real time: sit down each morning, brief Claude, generate, post. It works, but it wastes the system's biggest advantage.

Batch instead. Once a week, in one focused session: brief Claude for all seven days, generate the full set in Higgsfield, review the batch, schedule it. One afternoon buys a week — sometimes two. The rest of the week you are not making content, you are watching what performs and replying to comments.

Batching also makes you a better editor. Reviewing twenty images in one sitting, you see the persona as a whole — you catch the flat ones, the off-voice captions, the visual ruts. Day by day, you miss all of it.

Scaling: more surfaces, more avatars

Two ways to scale, and they compound.

More surfaces. One piece of content is not one post. A generated clip is an Instagram Reel, a TikTok, a YouTube Short, a Pinterest pin. Have Claude reformat the hook and caption per platform — same asset, four placements. This is the cheapest scaling there is; do it before anything else.

More avatars. Once one creator runs smoothly, the system is fully repeatable. A second avatar — different niche, different persona, same pipeline — is not double the work; the workflow is already built. Operators who treat this seriously end up running a small portfolio of AI creators, each a narrow niche, all on one back-end.

Do not start there. Get *one* avatar genuinely working — consistent, on-voice, with an audience that responds — before you clone the system. One channel that works beats a portfolio of mediocre ones.

How this actually makes money

Four honest paths, roughly easiest to hardest:

- **Your own product.** Highest margin, most control. The content builds the audience; a low-priced digital product — a guide, a template pack, a playbook — converts it. Content and product reinforce each other.
- **Affiliate.** Recommend tools and products the persona would genuinely use; earn on referrals. Works only if the recommendations are honest — the audience can tell.
- **Brand UGC.** Brands pay for content. An AI avatar already posting polished, on-brand content in a clear niche is a service you can sell — you produce their UGC with your pipeline.
- **Sponsorships.** Once an avatar has real reach, brands pay to appear. Last on the list because it requires the audience to already exist.

Across all four, the mechanism that converts is the same: give real value in the content, then make one clear, low-friction ask. A comment-to-DM offer — "comment a keyword, get the thing" — turns a viewer into a conversation, and a conversation into a sale, without the post ever feeling like an ad.

The honest part

What breaks: models update and a consistency setting drifts — re-check your Soul Character every few weeks. A platform changes its AI-content policy — keep labelling honestly and you are fine. An avatar's niche turns out not to monetize — kill it and reallocate; the system makes that cheap.

What to never do: do not pretend AI content is something it is not, do not spam the comment funnel, do not chase a niche you find boring. The system removes the *labour* of content. It does not remove the need to be worth following.

That is the playbook. A face that holds, a persona with a mind, an engine that writes and renders, and an operator — you — with taste and a cadence. The tools are commodity now. What you do with them is not. Go build something.

THE DISPATCH · FREE · MONTHLY

More like this, every morning.

The Flowi Daily Brief decodes one new AI release every day — what shipped, why it matters, what to do by Friday. Free monthly Dispatch arrives in your inbox.

useflowi.app/dispatch